

Aniruddha-Tapas /  
**Operating-Systems-Notes**<> **Code**

Issues

Pull requests 2

Actions

Projects

Security

Ir

**Operating-Systems-Notes / 5-Memory-Management.md****Aniruddha-Tapas** Update

7 years ago



319 lines (211 loc) · 7.14 KB

**Preview**

Code

Blame

Raw




# Memory Management

Operating systems:

- uses intelligently size containers
  - memory pages of segments
- Not all parts are needed at once
  - tasks operate on subset of memory
- Optimized for performance
  - reduce time to access state in memory
    - leads to better performance!

# Memory Management Goals



## Virtual vs Physical memory

- Allocate
  - allocation, replacement
- Arbitrate
  - address translation and validation


## Page-based Memory Management

- Allocate => pages => page frames
- Arbitrate => page tables

## Segment-based Memory Management

- Allocate => segments
- Arbitrate => segment registers

# Hardware Support



## Memory Management Unit (MMU)

- translate virtual to physical address
- reports faults (illegal access, permission, not present in memory)

## Registers

- pointers to page tables
- base and limit size, number of segments

## Cache

- Translation lookaside buffer
- Valid VA-PA translations using TLB

## Translation

- Actual PA generation done in hardware

# Page Tables

- OS creates page table per process
- On context switch, switch to valid page table
- Updates register that points to correct page table. E.g CR3 on x86 architecture

## Page Table Entry (PTE)



### Flags

- Present (valid/invalid)
- Dirty (written to)
- Accessed (for read or write)
- Protection bits => RWX

## Page Table Entry on x86



## Flags

- Present
- Dirty
- Accessed
- R/W permission bit 0: R only, 1: R/W
- U/S permission bit 0: usermode, 1: supervisor mode only
- others: caching related info (write through, caching disabled)
- unused: for future use


## Page faults

---

page Fault

## Page Table Size

---



- 32 bit architecture

- Page Table Entry (PTE) = 4 Bytes, including PFN + flags
- Virtual Page Number (VPN) =  $2^{32}/\text{page\_size}$
- Page size = 4KB (...8KB, 2MB, 4MB, 1GB)

Therefore Page Table Size =  $(2^{32} * 2^{12}) * 4B = 4MB$  (per process)

- for 64 bit architecture
  - Page Table Entry (PTE) = 8 Bytes
  - Page size = 4KB

Page Table Size =  $(2^{64} * 2^{12}) * 8B = 32PB$  (per process!)

- processes don't use entire address space
- even on 32 bit architecture, it will not always use all 4GB

But Page Table assumes an entry per VPN regardless, of whether corresponding virtual memory is needed or not.

## Hierarchical Page Tables

- outer page table or

On malloc, a new internal page table may be allocated.

**Address split:**

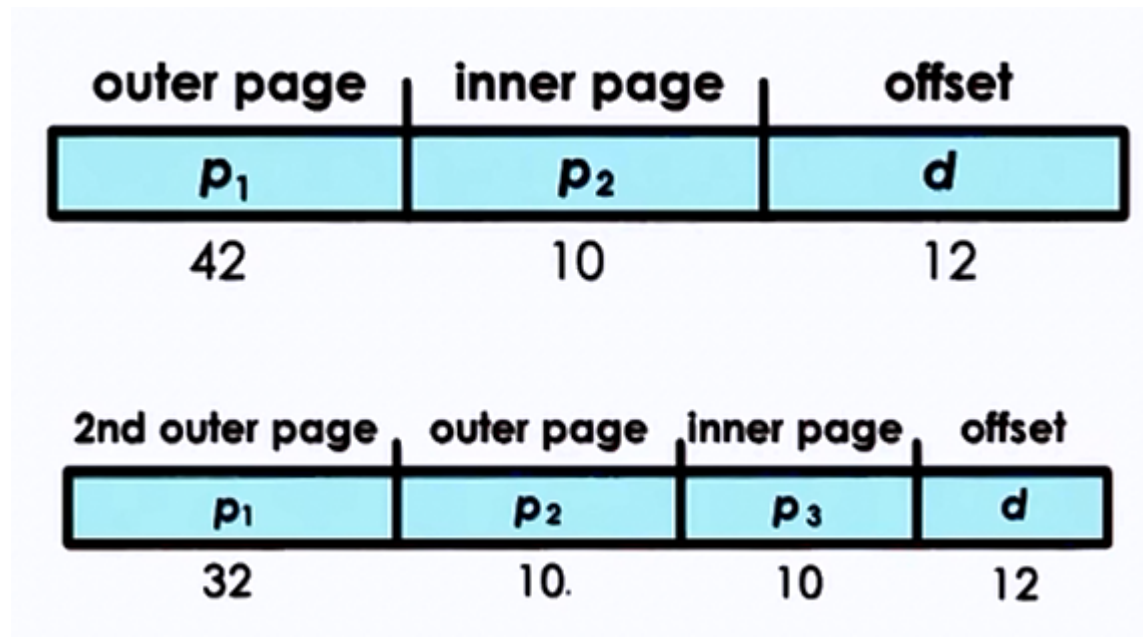
Page Number		offset
P1	P2	d
12	10	10

- inner table addresses  $\Rightarrow 2^{10} * \text{page\_size} = 2^{10} * 2^{10} = 1MB$

- don't need an inner table for each 1MB virtual memory gap

### Additional Layers

- page table directory pointer (3rd level)
- page table directory map (4th level)
- Important on 64 bit architectures
- larger and more sparse => larger gaps would save more internal page table components



## Tradeoffs of Multilevel Page Tables

### Advantages

- Smaller internal page tables/directories
- Granularity of coverage
  - Potentially reduced page table size

### Disadvantages

- More memory accesses required for translation
- increased translation latency

## Overheads of Address Translation


For each memory reference :

Single level page table	Four level page table
x1 access to PTE	x4 accesses to PTE

Single level page table	Four level page table
x1 access to mem	x1 access to mem

which results in slowdown.

## Page Table Cache

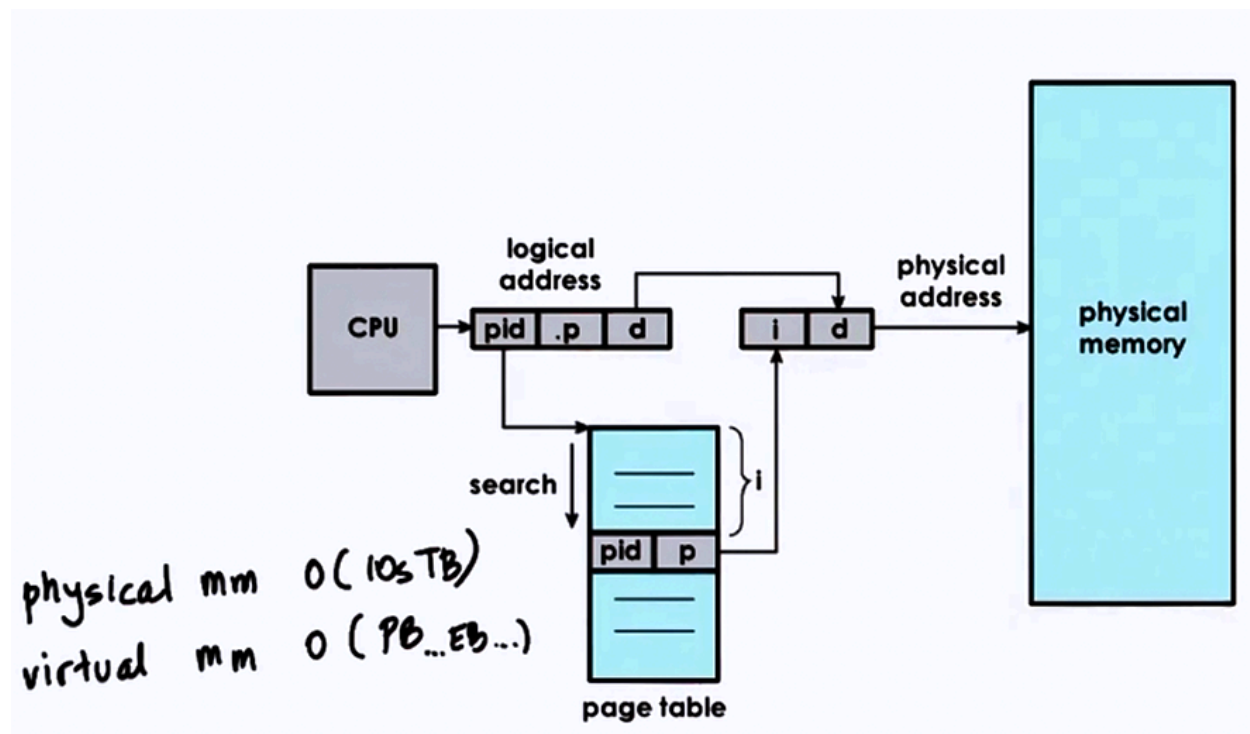


### Translation Lookaside Buffer


- MMU level address translation cache
- On TLB miss => page table access from memory
- has protection/validity bits
- small number of cached address => high TLB hit rate
  - temporal and spatial locality
- Example
  - x86 Core i7
    - per core : 64-entry data TLB  
128-entry instruction TLB
    - 512-entry shared second-level TLB



## Inverted Page Tables



## Hashing Page Tables




## Segmentation


Segmentation is the process of mapping virtual to physical memory using segments.

- Segments: arbitrary granularity (size)
  - e.g. code, heap, data, stack..
  - address = segment - selector + offset

- Segment
  - contiguous physical memory
  - segment size = segment base + limit registers



### Segmentation + Paging



## Page Size

- 10 bit offset => 1 KB page size [ $2^{10}$ ]
- 12 bit offset => 4 KB page size [ $2^{12}$ ]

In real world examples,

- Linux/x86 : 4 KB, 2MB, 1GB
- Solaris/Sparse: 8kB, 4MB, 2GB

	Large	Huge	page size	offset bits	reduction factor on page table size
	2 MB	1 GB	21 bits	30 bits	$512 \times 1024$

### Advantages

- larger pages
  - fewer page table entries, smaller page tables, more TLB hits

## Disadvantages

- internal fragmentation => wastes memory

## Memory Allocation

---

- Memory allocator
  - determines VA to PA mapping
  - address translation, page tables => simply determine PA from VA and check validity/permissions
- Kernel Level Allocators
  - kernel state, static process state
- User Level Allocators
  - dynamic process state (heap), malloc/free
  - e.g. d/malloc, jemalloc, Hoard, tcmalloc

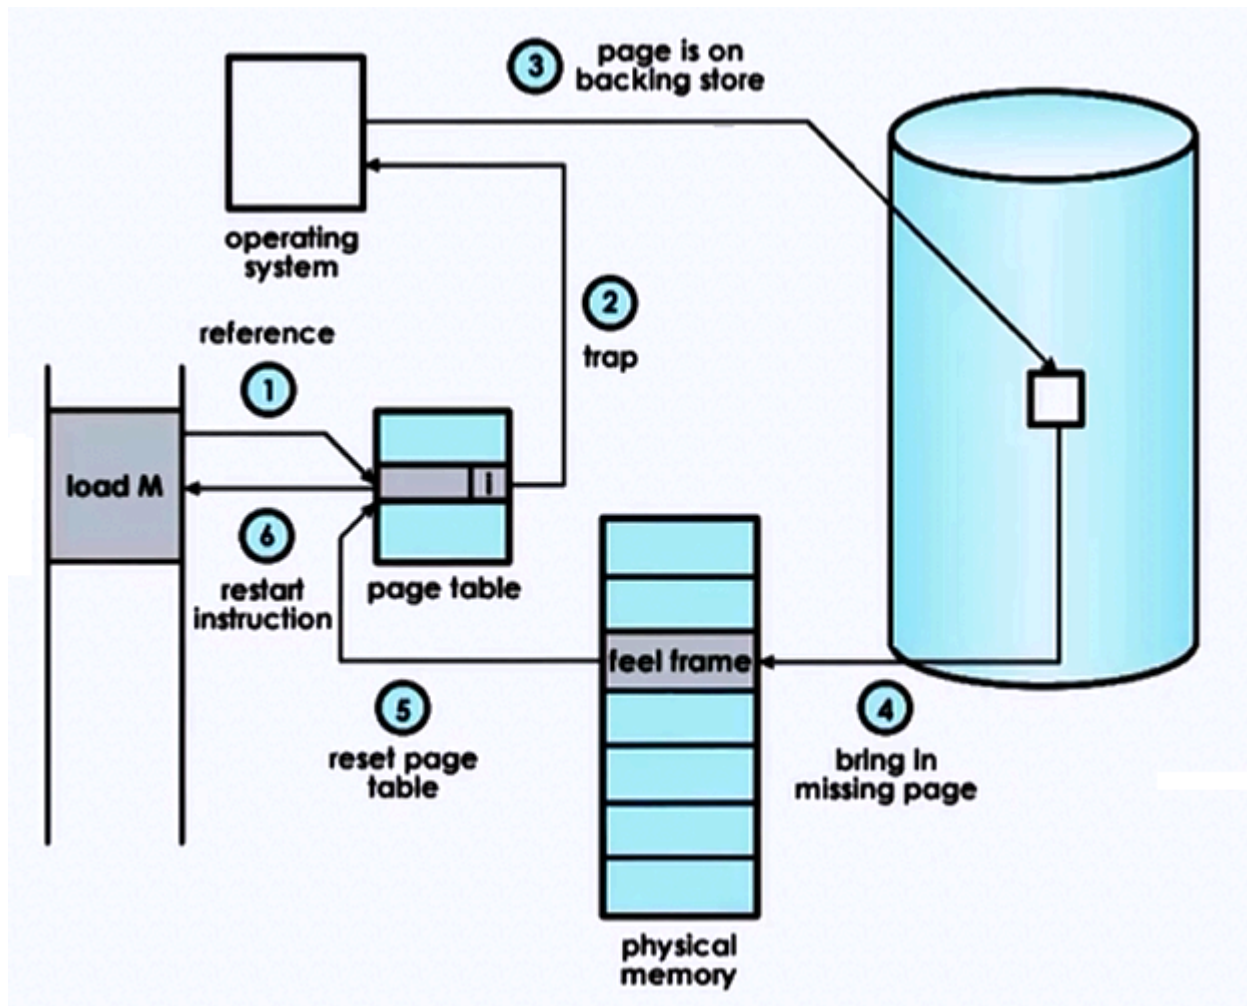
## Demand Paging

---

- Virtual Memory >> Physical Memory
  - virtual memory page is not always in physical memory
  - physical page frame saved and restored to/from secondary storage

### Demand paging:

- pages swapped in/out of memory & a swap partition (e.g. on a disk)



- Original PA != PA after swapping
  - if page is "pinned", swapping is disabled

### When pages should be swapped?

- page(out) daemon
- when memory usage is above threshold
- when CPU usage is below threshold

### Which page should be swapped out?

- pages that won't be used
- history based prediction
  - Least Recently Used (LRU policy). Access bit tracks if page is referenced.
- page that don't need to be written out
  - Dirty bit to track if modified
- avoid non-swappable pages

## Checkpointing

- Failure and Recovery management technique

- periodically save process state
- failure may be unavoidable but can restart from checkpoint, so recovery would be faster

### Simple Approach

- pause and save

### Better Approach

- write-protect and copy everything at once
- copy diffs of dirty pages for incremental checkpoints
  - rebuild from multiple diffs, or in background

Checkpointing can also be used in other services:

- Debugging
  - Rewind-Replay
  - rewind = restart from checkpoint
  - gradually go back to earlier checkpoints until error is found
- Migration
  - continue on another machine
  - disaster recovery
  - consolidation
  - repeated checkpoints in a fast loop until pause and copy becomes acceptable (or unavoidable)