



Operating-Systems-Notes / 7-Synchronization.md



Aniruddha-Tapas Commit

7 years ago



132 lines (95 loc) · 3 KB

Preview

Code

Blame

Raw



Synchronization

Waiting for other processes, so that they can continue working together

- may repeatedly check to continue
 - sync using spinlocks
- may wait for a signal to continue
 - sync using mutexes and condition variables
- waiting hurts performance
 - CPUs waste cycles for checking; cache effects

Limitation of mutexes and condition variables

- Error prone/correctness/ease of use
 - unlock wrong mutex, signal wrong condition variable
- Lack of expressive power
 - helper variables for access or priority control

Low-level support: hardware atomic instructions

Synchronization constructs

1. Spinlocks (basic sync construct)
 - Spinlock is like a mutex
 - mutual exclusion
 - lock and unlock(free)
 - but, lock == busy => spinning

2. Semaphores

- common sync construct in OS kernels
- like a traffic light: Stop and Go
- like mutex, but more general

Semaphore == integer value

- on init
 - assigned a max value (positive int) => max count
- on try(wait)
 - if non-zero, decrement and proceed => counting semaphore
- if initialized with 1
 - semaphore == mutex(binary semaphore)
- on exit(post)
 - increment

Syncing different types of accesses

Reader/Writer locks

read (don't modify)	write (always modify)
shared access	exclusive access

- RW locks
 - specify type of access, then lock behaves accordingly

Monitors (highlevel construct)

- shared resource
- entry resource
- possible condition variables
- On entry:
 - lock, check
- On exit:
 - unlock, check, signal

More synchroniaztion constructs

- serializers
- path expressions
- barriers

- rendezvous points
- optimistic wait-free sync (RCU) [Read Copy Update]

All need hardware support.

Need for hardware support

- Problem
 - concurrent check/update on different CPUs can overlap

Atomic instructions

Critical section with hardware supported synchronization

Hardware specific

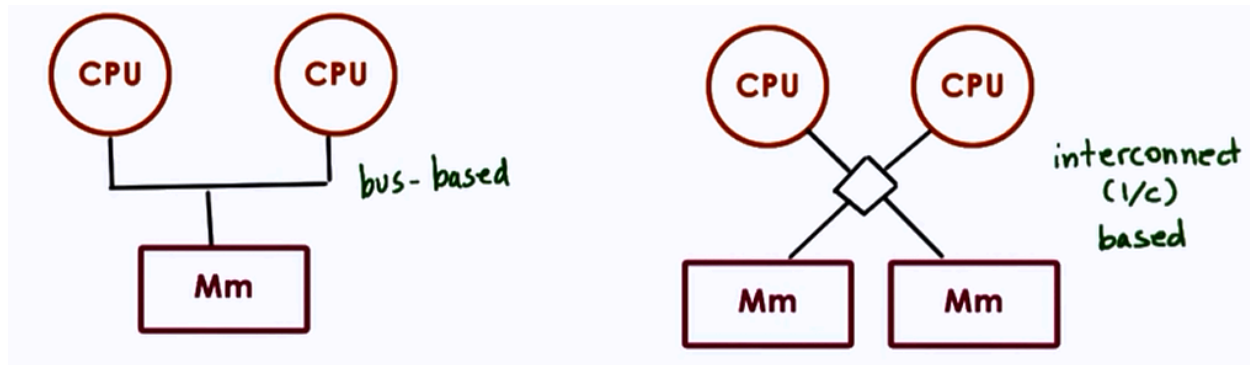
- test-and-set
 - returns(tests) original values and sets new-value!= 1 (busy) automatically
 - first thread: test-and-set(lock) => 0 : free
 - next ones: test-and-set(lock) => 1 busy
 - reset lock to 1, but that's okay
 - + : Latency
 - + : minimal (Atomic)
 - + : Delay potentially min
 - - : Contention processors go to memory on each spin - To reduce contention, introduce delay - Static(based on a fixed value) or Dynamic(backoff based, random delay)
- read-and-increment
- compare-and-swap

Guarantees

- atomicity
- mutual exclusion
- queue all concurrent instructions but one

Shared Memory Multiprocessors

Also called symmetric multiprocessors (SMP)



- Caches
 - hide memory latency, "memory" further away due to contention
 - no-write, write-through, write-back

Cache Coherence

