



## Operating-Systems-Notes / 9-Virtualization.md



Aniruddha-Tapas Update

7 years ago



212 lines (149 loc) · 6.03 KB

Preview

Code


Blame

Raw



# Virtualization

Virtualization allows concurrent execution of multiple OSs and their applications on the same physical machine.



- Virtual resources : each OS thinks that it "owns" hardware resources
- Virtual machine (VM) : OS + applications + virtual resources (guest domain)
- Virtualization layer : management of physical hardware (virtual machine monitor, hypervisor)

## Defining Virtual Machine

A Virtual Machine is an efficient, isolated duplicate of the real machine.

- Supported by a Virtual Machine Monitor (VMM):
  - i. provides environment essentially identical with the original machine
  - ii. programs show only minor decrease in speed at worst
  - iii. VMM is in complete control of the system resources

## VMM goals

---

- Fidelity
- Performance
- Safety and Isolation

## Virtualization advantages


---

- consolidation
  - decrease cost, improve manageability
- migration
  - availability, reliability
- security, debugging, support for legacy OS

## Two main Virtualization Models:

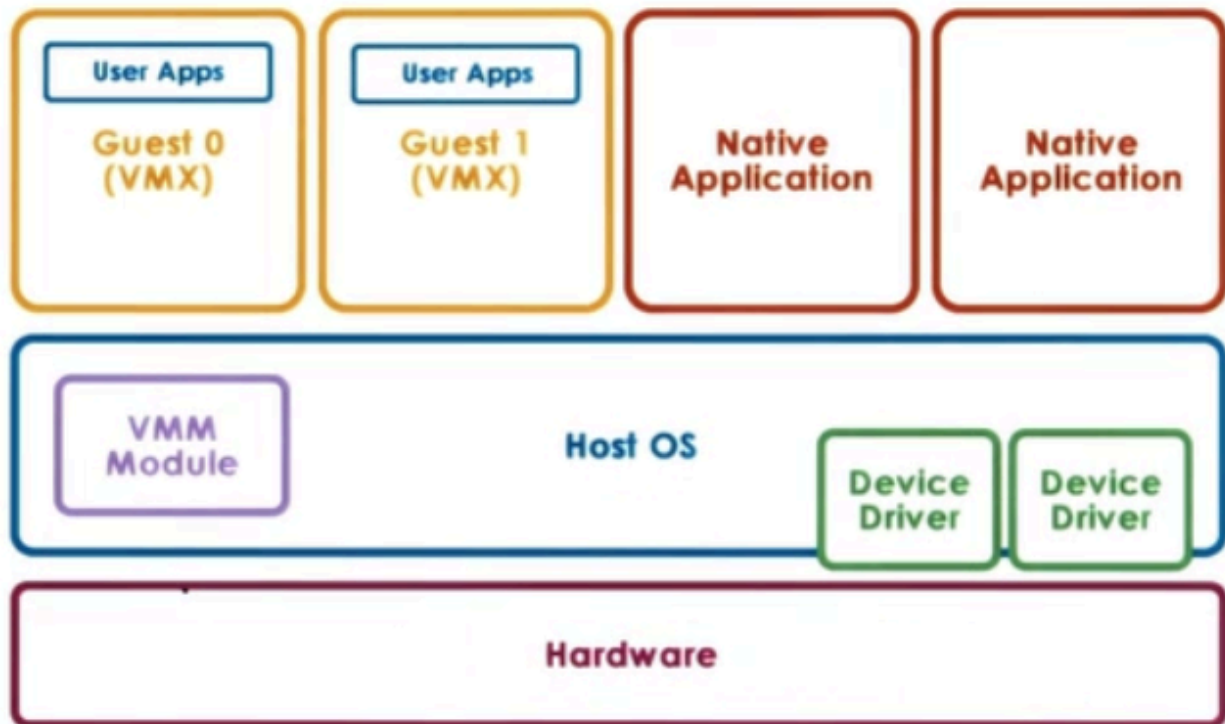
---

### 1. Bare-metal or Hypervisor based (Type 1)



- VMM (hypervisor) manages all hardware resources and supports execution of VMs
- privileged, secure VM to deal with devices (and other configuration and management tasks)
- Adopted by Xen (OpenSource or Citrix Xen Server) and ESX (VMware)

## 1. Hosted (Type 2)



- Host owns all hardware
- Special VMM module provides hardware interfaces to VMs and deals with VM context switching

## Virtualization requirements

- Present virtual platform interface to VMs
  - virtualize CPU, memory, devices
- Provide isolation across VMs
  - preemption, MMU for address translation and validation
- Protect guest OS from applications
  - can't run guest OS and applications at same protection level
- Protect VMs from guest OS
  - can't run guest OS and VMMs at same protection level

## Hardware protection levels

Commodity hardware has more than two protection levels



- x86 has 4 protection levels (rings)
  - ring 3 : lowest privilege (applications)
  - ring 1 : OS
  - ring 0 : highest privilege (hypervisor)
- and 2 protection modes
  - non root : VMs
    - ring 3 : apps
    - ring 0 : OS
  - root :
    - ring 0 : hypervisor

## Process Virtualization (Trap-and-Emulate)

- Guest instruments
  - executed directly by hardware
  - for non-privileged operations : hardware speeds => efficiency
  - for privileged operations : trap to hypervisor
- Hypervisor determines what needs to be done:
  - if illegal operation : terminate VM
  - if legal operation : emulate the behaviour the guest OS was expecting from the hardware

## Problems with Trap-and-Emulate

---

- 17 privileged information do not trap but fail silently
- Hypervisor doesn't know, so it doesn't try to change settings
- OS doesn't know, so assumes change was successful

## Binary Translation

---

**Goal** : Full Virtualization i.e. guest OS is not modified

**Approach** : Dynamic Binary Translation

1. Inspect code blocks to be executed
2. If needed, translate to alternate instruction sequence
  - e.g. to emulate desired behaviour, possibly avoid traps
3. Otherwise run at hardware speeds
  - cache translated blocks to amortize translation costs

## Paravirtualization

---

**Goal** : Performance; give up on modified guest OSs

**Approach** : Paravirtualization : modify guest OSs so that

- it knows it is running virtualized
- it makes explicit calls to hypervisor (hypercalls)
- hypercalls (~ system calls)
  - package context information
  - specify desired hypercall
  - trap to VMM
- Xen : open-source hypervisor

## Memory virtualization

---

- Full virtualization
  - all guests expect contiguous physical memory starting at 0
  - virtual vs physical vs machine addresses and page frame numbers
  - still leverages hardware (MMU, TLB..)
- Option 1
  - guest page table : VA => PA
  - hypervisor : PA => MA

- too expensive!
- Option 2
  - guest page tables : VA => PA
  - hypervisor shadow PT : VA => MA
  - hypervisor maintains consistence
    - e.g. invalidate on context switch, write protect guest PT to track new mappings
- Paravirtualized
  - guest aware of virtualization
  - no longer strict requirement on contiguous physical memory starting at 0
  - explicitly registers page tables with hypervisor
  - can "batch" page tables updates to reduce VM exits
  - other optimazations

Overheads eliminated or reduced on newer platforms

## Device Virtualization


---

- For CPUs and Memory
  - less diversity, Intruction-Set-Architecture(ISA) level
  - Standardization of interface
- For Devices
  - high diversity
  - lack of standard specification of device interface and behaviour

### 3 key models for Device Virtualization:

#### 1. Pass through model

Approach: VMM-level-driver configures device access permissions



### Advantages

- VM provided with exclusive and direct (VMM bypass) access to the device


### Disadvantages

- Device sharing difficult
- VMM must have exact type of device as what VM expects
- VM migration tricky

## 2. Hypervisor - Direct model

Approach:

- VMM interrupts all device accesses
- Emulate device operations
  - translate to generic I/O operations
  - traverse VMM-resident I/O stack
  - invoke VMM-resident driver



### Advantages

- VM decoupled from physical device
- Sharing, migration, dealing with device specifics

### Disadvantages

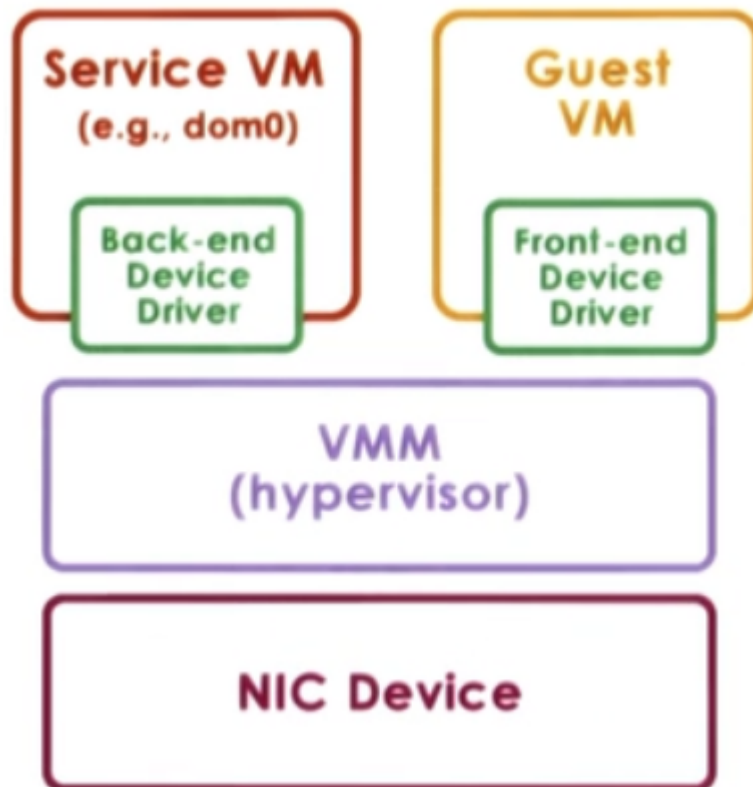
- Latency of device operations
- Device driver ecosystem complexities in Hypervisor

## 3. Split Device-Driver model

Approach:

- Device access control split between
- Emulate device operations
  - front-end driver in guest VM (device API)
  - back-end driver in service VM (or Host)
  - modified guest drivers
    - i.e. limited to paravirtualized guests





### Advantages

- Eliminate emulation overhead
- Allow for better management of shared devices