



117 lines (94 loc) · 4.2 KB

Preview

Code

Blame

Raw





Distributed File Systems

- Accessed via well defined interface
 - access via Virtual File Systems
- Focus on consistent state
 - tracking state, file update, cache coherence
- Mixed distribution models possible
 - replicates vs partitioned, peer-like systems

DFS models

- Client Server on different machines
- File server distributed on multiple machines
 - replicated (each server : all files)
 - partitioned (each server : parts of files)
 - both (files partitioned, each partition replicates)
- Files stored on and served from all machines (peers)
 - blurred distinction between clients and servers

Remote File Service : Extremes



1. Extreme1 : Upload/Download
 - like FTP, SVN
 - + local read/writes at client
 - - entire file download/upload evn for small accesses
 - - server gives up contro;
2. Extreme2 : True Remote File Access
 - Every access to remote file, nothing done locally
 - + file access centralized, easy to reason about consistency
 - - every file operation pays network cost, limits server scalability

Remote File Service : A compromise

A more practical Remote File access (with Caching)

1. Allow clients to store parts of files locally (blocks)
 - + low latency on file operations
 - + server load reduces => more scalable
2. Force clients to interact with server (frequently)
 - + server has insights into what clients are doing
 - + server has control into which accesses can be permitted => easier to maintain consistency
 - - server more complex, requires different file sharing semantics

Stateless vs Stateful File server

Stateless	Stateful
Keeps no state; Okay with extreme models, but can't support 'practical' model	Keeps client state needed for 'practical' model to track what is cached/accessed
- Can't support caching and consistency management	+ Can support locking, caching, incremental operations
- Every request self-contained. => more bits transferred	- Overheads to maintain state and consistency. Depends on caching mechanism and consistency protocol.
+ No resources are used on server side (CPU, MM). On failure just restart	- On failure, need checkpointing and recovery mechanisms

Caching state in a DFS

- Locally clients maintain portion of state (e.g. file blocks)
- Locally clients perform operations on cached state (e.g. open/read/write)
- requires coherent mechanisms



System	How	When
SMP	Write-update/Write-invalidate	On write
DFS	Client/Server-driven	On demand, periodically, on open..

- Files or File blocks can be (with 1 server and multiple clients) cached in:
 - in client memory
 - on client storage device (HDD/SDD)
 - in buffer cache in memory on server
 - (usefulness will depend on client load, request interleaving)
- File Sharing Semantics in DFS
- Session semantics (between open-close => Session)
 - write-back on close(), update on open()
 - easy to reason, but may be insufficient
- Periodic updates
 - client writes-back periodically
 - clients have a "lease" on cached data (not exclusively necessary)
 - servers invalidates periodically => provides biunds on "inconsistency"
 - augment with flush()/sync() API
- Immutable files => never modify, new files created
- Transactions => all changes atomic

Replication vs Partitioning

	Replication	Partitioning
	Each machine holds all files	Each machine has subset of files
Advantages	Load balancing, availability, fault tolerance	Availability vs single server DFS; Scalability with file system size; single file writes simpler
Disadvantages	Write becomes more complex <ul style="list-style-type: none">- Synchronous to all- or, write to one, then	On failure, lose portion of data load balancing harder, if not balanced, then hot-spots possible

propagate to others replicas must be reconciled e.g. Voting

- Can combine both techniques
 - Replicate each partition!