

Started on	Saturday, 16 March 2024, 1:32 PM
State	Finished
Completed on	Saturday, 16 March 2024, 3:45 PM
Time taken	2 hours 12 mins
Grade	11.98 out of 15.00 (79.88%)

Question 1

Partially correct

Mark 0.31 out of 0.50

It is proposed that when a process does an illegal memory access, xv6 terminate the process by printing the error message "Illegal Memory Access". Select all the changes that need to be done to xv6 for this as True (Note that the changes proposed here may not cover the exhaustive list of all changes required) and the un-necessary/wrong changes as False.

Required	Un-necessary/Wrong		
<input type="radio"/>	<input checked="" type="radio"/>	Add code that checks if the illegal memory access trap was due to an actual illegal memory access.	
<input checked="" type="radio"/>	<input type="radio"/>	Change in the Makefile and instruct cc/ld to start the code of each program at some address other than 0	
<input checked="" type="radio"/>	<input type="radio"/>	Change exec to treat text/data sections separately and call allocuvn() with proper flags for page table entries	
<input checked="" type="radio"/>	<input type="radio"/>	Ensure that the address 0 is mapped to invalid	
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Change mappages() to set specified permissions on each page table entry	
<input type="radio"/>	<input checked="" type="radio"/>	Mark each page as readonly in the page table mappings	
<input checked="" type="radio"/>	<input type="radio"/>	Change allocuvn() to call mappages() with proper permissions on each page table entry	
<input checked="" type="radio"/>	<input type="radio"/>	Handle the Illegal memory acceess trap in trap() function, and terminate the currently running process.	

Add code that checks if the illegal memory access trap was due to an actual illegal memory access.: Un-necessary/Wrong

Change in the Makefile and instruct cc/ld to start the code of each program at some address other than 0: Required

Change exec to treat text/data sections separately and call allocuvn() with proper flags for page table entries: Required

Ensure that the address 0 is mapped to invalid: Required

Change mappages() to set specified permissions on each page table entry: Un-necessary/Wrong

Mark each page as readonly in the page table mappings: Un-necessary/Wrong

Change allocuvn() to call mappages() with proper permissions on each page table entry: Required

Handle the Illegal memory acceess trap in trap() function, and terminate the currently running process.: Required

## Question 2

Correct

Mark 0.50 out of 0.50

Given below are statements about concurrency and parallelism

Select T/F

A concurrent system can allow more than one task to progress, whereas a parallel system can perform more than one task at the same time.

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	Both concurrency and parallelism are the same.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A concurrent system can allow more than one task to progress, whereas a parallel system can perform more than one task at the same time.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Parallel systems allow more than one task to progress while concurrent systems do not.	✓
<input type="radio"/>	<input checked="" type="radio"/>	It is not possible to have concurrency without parallelism.	✓
<input checked="" type="radio"/>	<input type="radio"/>	It is possible to have concurrency without parallelism	✓
<input type="radio"/>	<input checked="" type="radio"/>	It is possible to have parallelism without concurrency	✓
<input type="radio"/>	<input checked="" type="radio"/>	A concurrent system allows more than one task to progress while a parallel system does not.	✓

Both concurrency and parallelism are the same.: False

A concurrent system can allow more than one task to progress, whereas a parallel system can perform more than one task at the same time.: True

Parallel systems allow more than one task to progress while concurrent systems do not.: False

It is not possible to have concurrency without parallelism.: False

It is possible to have concurrency without parallelism: True

It is possible to have parallelism without concurrency: False

A concurrent system allows more than one task to progress while a parallel system does not.: False

Question **3**

Partially correct

Mark 0.19 out of 0.50

Suppose a file is to be created in an ext2 file system, in an existing directory */a/b/*. Select from below, the list of blocks that may need modification.

Select one or more:

- ☐ a. link count on */a/b/* inode
- ☒ b. inode bitmap referring to */a/b/* ✖
- ☐ c. data blocks of */a/*
- ☒ d. superblock ✔
- ☐ e. inode bitmap in some block group
- ☐ f. existing data blocks of */a/b/*
- ☒ g. new data block in some block group ✔
- ☐ h. group descriptor(s)
- ☒ i. inode of */a/b/* ✔
- ☒ j. block bitmap in some block group ✔
- ☐ k. inode of */a/*
- ☒ l. inode table in some block group ✔

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: superblock, group descriptor(s), inode of */a/b/*, existing data blocks of */a/b/*, inode table in some block group, inode bitmap in some block group, block bitmap in some block group, new data block in some block group

Question **4**

Correct

Mark 0.50 out of 0.50

The "push 0" in vectors.S is

- ☐ a. To be filled in as the return value of the system call
- ☐ b. To indicate that it's a system call and not a hardware interrupt
- ☐ c. A placeholder to match the size of struct trapframe
- ☒ d. Place for the error number value ✔

The correct answer is: Place for the error number value

Question **5**

Partially correct

Mark 0.42 out of 0.50

Select the correct statements about hard and soft links

Select one or more:

- ☒ a. Soft links can span across partitions while hard links can't ✓
- ☐ b. Deleting a soft link deletes both the link and the actual file
- ☐ c. Soft links increase the link count of the actual file inode
- ☒ d. Hard links share the inode ✓
- ☐ e. Deleting a soft link deletes only the actual file
- ☐ f. Hard links enforce separation of filename from it's metadata in on-disk data structures.
- ☒ g. Deleting a hard link deletes the file, only if link count was 1 ✓
- ☐ h. Deleting a hard link always deletes the file
- ☒ i. Hard links increase the link count of the actual file inode ✓
- ☐ j. Hard links can span across partitions while soft links can't
- ☒ k. Deleting a soft link deletes the link, not the actual file ✓
- ☐ l. Soft link shares the inode of actual file

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: Soft links can span across partitions while hard links can't, Hard links increase the link count of the actual file inode, Deleting a soft link deletes the link, not the actual file, Deleting a hard link deletes the file, only if link count was 1, Hard links share the inode, Hard links enforce separation of filename from it's metadata in on-disk data structures.

Question **6**

Correct

Mark 0.50 out of 0.50

Select the statement that most correctly describes what setupkvm() does

- ☐ a. creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array and makes kpgdir point to it
- ☐ b. creates a 2-level page table for the use of the kernel, as specified in gdt\_desc
- ☐ c. creates a 1-level page table for the use by the kernel, as specified in kmap[] global array
- ☒ d. creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array ✓

The correct answer is: creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array

## Question 7

Correct

Mark 0.50 out of 0.50

Mark the statements as True/False, with respect to the use of the variable "chan" in struct proc.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	chan stores the address of the variable, representing a condition, for which the process is waiting.	✓
<input checked="" type="radio"/>	<input type="radio"/>	The value of 'chan' is changed only in sleep()	✓
<input checked="" type="radio"/>	<input type="radio"/>	in xv6, the address of an appropriate variable is used as a "condition" for a waiting process.	✓
<input type="radio"/>	<input checked="" type="radio"/>	when chan is NULL, the 'state' in proc must be RUNNABLE.	✓
<input checked="" type="radio"/>	<input type="radio"/>	When chan is not NULL, the 'state' in struct proc must be SLEPING	✓
<input type="radio"/>	<input checked="" type="radio"/>	chan is the head pointer to a linked list of processes, waiting for a particular event to occur	✓
<input type="radio"/>	<input checked="" type="radio"/>	Changing the state of a process automatically changes the value of 'chan'	✓
<input checked="" type="radio"/>	<input type="radio"/>	'chan' is used only by the sleep() and wakeup1() functions.	✓

chan stores the address of the variable, representing a condition, for which the process is waiting.: True

The value of 'chan' is changed only in sleep(): True

in xv6, the address of an appropriate variable is used as a "condition" for a waiting process.: True

when chan is NULL, the 'state' in proc must be RUNNABLE.: False

When chan is not NULL, the 'state' in struct proc must be SLEPING: True

chan is the head pointer to a linked list of processes, waiting for a particular event to occur: False

Changing the state of a process automatically changes the value of 'chan': False

'chan' is used only by the sleep() and wakeup1() functions.: True

Consider this program.

Some statements are identified using the // comment at the end.

Assume that = is an atomic operation.

```
#include <stdio.h>
#include <pthread.h>
long c = 0, c1 = 0, c2 = 0, run = 1;
void *thread1(void *arg) {
    while(run == 1) { //E
        c = 10; //A
        c1 = c2 + 5; //B
    }
}
void *thread2(void *arg) {
    while(run == 1) { //F
        c = 20; //C
        c2 = c1 + 3; //D
    }
}
int main() {
    pthread_t th1, th2;
    pthread_create(&th1, NULL, thread1, NULL);
    pthread_create(&th2, NULL, thread2, NULL);
    sleep(2);
    run = 0;
    fprintf(stdout, "c = %ld c1+c2 = %ld c1 = %ld c2 = %ld \n", c, c1+c2, c1, c2);
    fflush(stdout);
}
```

Which statements are part of the critical Section?

Yes	No		
<input checked="" type="radio"/>	<input type="radio"/>	D	✓
<input type="radio"/>	<input checked="" type="radio"/>	A	✗
<input type="radio"/>	<input checked="" type="radio"/>	C	✓
<input checked="" type="radio"/>	<input type="radio"/>	B	✓
<input type="radio"/>	<input checked="" type="radio"/>	E	✓
<input type="radio"/>	<input checked="" type="radio"/>	F	✓

- D: Yes
- A: No
- C: No
- B: Yes
- E: No
- F: No

## Question 9

Partially correct

Mark 0.33 out of 0.50

Mark statements about deadlocks as True or false

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	A deadlock must involve at least two processes	✗
<input type="radio"/>	<input checked="" type="radio"/>	Deadlocks are not possible if there is no race	✗
<input checked="" type="radio"/>	<input type="radio"/>	A deadlock necessarily requires a cycle in the resource allocation graph	✓
<input checked="" type="radio"/>	<input type="radio"/>	A deadlock is possible only if all the 4 conditions of mutual exclusion, cyclic wait, hold and wait, and no preemption are satisfied	✓
<input type="radio"/>	<input checked="" type="radio"/>	Deadlocks are the same as livelocks	✓
<input checked="" type="radio"/>	<input type="radio"/>	Cycle in the resource allocation graph does not necessarily mean a deadlock	✓

A deadlock must involve at least two processes: False

Deadlocks are not possible if there is no race: True

A deadlock necessarily requires a cycle in the resource allocation graph: True

A deadlock is possible only if all the 4 conditions of mutual exclusion, cyclic wait, hold and wait, and no preemption are satisfied: True

Deadlocks are the same as livelocks: False

Cycle in the resource allocation graph does not necessarily mean a deadlock: True

Question **10**

Correct

Mark 0.50 out of 0.50

Match the code with it's functionality

S = 5

Wait(S)

Critical Section

Signal(S)

Counting semaphore



S = 0

P1:

Statement1;

Signal(S)

Execution order P1, then P2



P2:

Wait(S)

Statment2;

S1 = 0; S2 = 0;

P2:

Statement1;

Signal(S2);

P1:

Wait(S2);

Statemetn2;

Signal(S1);

Execution order P2, P1, P3



P3:

Wait(S1);

Statement S3;

S = 1

Wait(S)

Critical Section

Signal(S);

Binary Semaphore for mutual exclusion



Your answer is correct.

The correct answer is: S = 5

Wait(S)

Critical Section

Signal(S) → Counting semaphore, S = 0

P1:

Statement1;

Signal(S)

P2:

Wait(S)

Statment2; → Execution order P1, then P2, S1 = 0; S2 = 0;

P2:

Statement1;

Signal(S2);

P1:

Wait(S2);

Statemetn2;

Signal(S1);

P3:

Wait(S1);

Statement S3; → Execution order P2, P1, P3, S = 1

Wait(S)

Critical Section

Signal(S); → Binary Semaphore for mutual exclusion



Question **11**

Correct

Mark 0.50 out of 0.50

Which of the following is DONE by allocproc() ?

- ☐ a. setup the contents of the trapframe of the process properly
- ☒ b. allocate PID to the process ✓
- ☒ c. allocate kernel stack for the process ✓
- ☐ d. setup kernel memory mappings for the process
- ☒ e. setup the trapframe and context pointers appropriately ✓
- ☐ f. ensure that the process starts in trapret()
- ☒ g. Select an UNUSED struct proc for use ✓
- ☒ h. ensure that the process starts in forkret() ✓

The correct answers are: Select an UNUSED struct proc for use, allocate PID to the process, allocate kernel stack for the process, setup the trapframe and context pointers appropriately, ensure that the process starts in forkret()

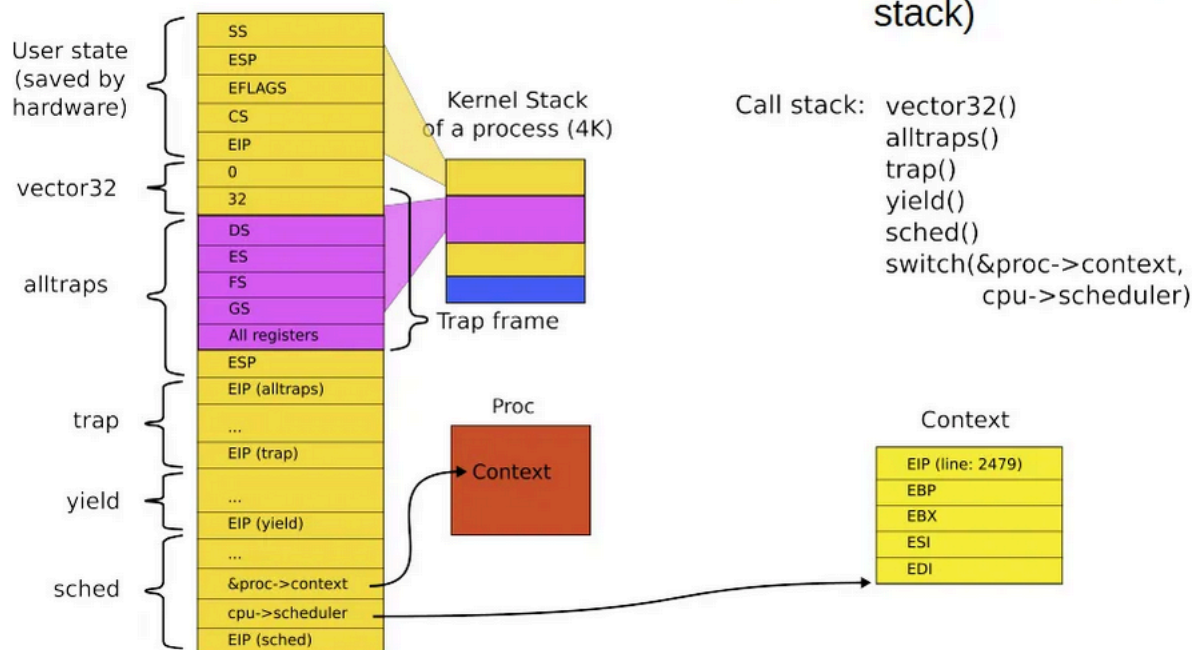
## Question 12

Correct

Mark 0.50 out of 0.50

Mark statements as True/False, w.r.t. the given diagram

## Stack inside switch() and its two arguments (passed on the stack)



True	False		
<input type="radio"/>	<input checked="" type="radio"/>	The "context" yellow coloured box, pointed to by cpu->scheduler is on the kernel stack of the scheduler.	✓
<input type="radio"/>	<input checked="" type="radio"/>	This is a diagram of switch() called from scheduler()	✓ No. diagram of switch() called from sched()
<input checked="" type="radio"/>	<input type="radio"/>	The "ESP" (second entry from top) is stack pointer of user-stack of process, while the "ESP" (first entry below pink region) is the trapframe pointer on kernel stack of process.	✓
<input checked="" type="radio"/>	<input type="radio"/>	The blue shaded part in "kernel stack of a process(4k)" refers to remaining part of stack (not used yet)	✓
<input type="radio"/>	<input checked="" type="radio"/>	The diagram is wrong because it shows the user stack and kernel stack together (continuous), but in practice they are separate	✓ diagram shows only kernel stack
<input checked="" type="radio"/>	<input type="radio"/>	The diagram is correct	✓

The "context" yellow coloured box, pointed to by cpu-&gt;scheduler is on the kernel stack of the scheduler.: False

This is a diagram of switch() called from scheduler(): False

The "ESP" (second entry from top) is stack pointer of user-stack of process, while the "ESP" (first entry below pink region) is the trapframe pointer on kernel stack of process.: True

The blue shaded part in "kernel stack of a process(4k)" refers to remaining part of stack (not used yet): True

The diagram is wrong because it shows the user stack and kernel stack together (continuous), but in practice they are separate: False

The diagram is correct: True

Question **13**

Correct

Mark 0.50 out of 0.50

Select the correct statements about paging (not demand paging) mechanism

Select one or more:

- ☐ a. An invalid entry on a page means, either it was illegal memory reference or the page was not present in memory.
- ☒ b. An invalid entry on a page means, it was an illegal memory reference ✓
- ☐ c. User process can update it's own PTBR
- ☒ d. The PTBR is loaded by the OS ✓
- ☐ e. User process can update it's own page table entries
- ☒ f. Page table is accessed by the MMU as part of execution of an instruction ✓
- ☐ g. Page table is accessed by the OS as part of execution of an instruction
- ☒ h. OS creates the page table for every process ✓

Your answer is correct.

The correct answers are: OS creates the page table for every process, The PTBR is loaded by the OS, Page table is accessed by the MMU as part of execution of an instruction, An invalid entry on a page means, it was an illegal memory reference

Question 14

Partially correct

Mark 0.13 out of 0.50

Map ext2 data structure features with their purpose

Used  
directories  
count in  
group  
descriptor

Choose...

Combining  
file type  
and  
access  
rights in  
one  
variable

saves 1 byte of space



Inode table  
location in  
Group  
Descriptor

Choose...

rec\_len  
field in  
directory  
entry

limits total number of files that can belong to a group



File Name  
is padded

aligns all memory accesses on word boundary, improving performance



Inode  
bitmap is  
one block

allows holes and linking of entries in directory



Free  
blocks  
count in  
superblock  
and group  
descriptor

Choose...

Mount  
count in  
superblock

to enforce file check after certain amount of mounts at boot time



A group

All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to



Many  
copies of  
Superblock

Choose...

Block  
bitmap is  
one block

Choose...

Inode table

Obvious, as it's per group and not per file-system



Your answer is partially correct.

You have correctly selected 3.

The correct answer is: **Used directories count in group descriptor** → attempt is made to evenly spread the first-level directories, this count is used there, **Combining file type and access rights in one variable** → saves 1 byte of space, **Inode table location in Group Descriptor** → Obvious, as it's per group and not per file-system, **rec\_len field in directory entry** → allows holes and linking of entries in directory, File Name is padded → aligns all memory accesses on word boundary, improving performance, **Inode bitmap is one block** → limits total number of files that can belong to a group, **Free blocks count in superblock and group descriptor** → Redundancy to help fsck restore consistency, **Mount count in superblock** → to enforce file check after certain amount of mounts at boot time, **A group** → Try to keep all the data of a directory and it's file close together in a group, **Many copies of Superblock** → Redundancy to ensure the most crucial data structure is not lost, **Block bitmap is one block** → Limits the size of a block group, thus improvising on purpose of a group, **Inode table** → All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk

Question **15**

Partially correct

Mark 0.25 out of 0.50

Match pairs

peterson	per process flag, global turn variable	✓
semaphore	wait() and signal()	✓
spinlock	lock() and unlock()	✗
mutex	atomic test and set with loop	✗

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: peterson → per process flag, global turn variable, semaphore → wait() and signal(), spinlock → atomic test and set with loop, mutex → lock() and unlock()

Question 16

Correct

Mark 0.50 out of 0.50

Consider the image given below, which explains how paging works.

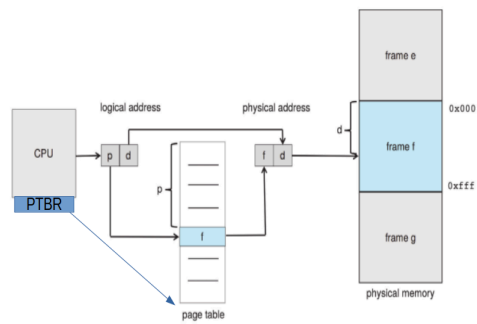


Figure 9.8 Paging hardware.

Mention whether each statement is True or False, with respect to this image.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Maximum Size of page table is determined by number of bits used for page number	✓
<input type="radio"/>	<input checked="" type="radio"/>	The page table is indexed using frame number	✓
<input checked="" type="radio"/>	<input type="radio"/>	The PTBR is present in the CPU as a register	✓
<input checked="" type="radio"/>	<input type="radio"/>	The page table is indexed using page number	✓
<input checked="" type="radio"/>	<input type="radio"/>	The physical address may not be of the same size (in bits) as the logical address	✓
<input type="radio"/>	<input checked="" type="radio"/>	Size of page table is always determined by the size of RAM	✓
<input checked="" type="radio"/>	<input type="radio"/>	The page table is itself present in Physical memory	✓
<input type="radio"/>	<input checked="" type="radio"/>	The locating of the page table using PTBR also involves paging translation	✓

Maximum Size of page table is determined by number of bits used for page number: True  
The page table is indexed using frame number: False  
The PTBR is present in the CPU as a register: True  
The page table is indexed using page number: True  
The physical address may not be of the same size (in bits) as the logical address: True  
Size of page table is always determined by the size of RAM: False  
The page table is itself present in Physical memory: True  
The locating of the page table using PTBR also involves paging translation: False

## Question 17

Correct

Mark 0.50 out of 0.50

Select the correct statements about sched() and scheduler() in xv6 code

- ☒ a. scheduler() switches to the selected process's context ✓
- ☒ b. When either sched() or scheduler() is called, it results in a context switch ✓
- ☒ c. sched() switches to the scheduler's context ✓
- ☒ d. Each call to sched() or scheduler() involves change of one stack inside swtch() ✓
- ☒ e. When either sched() or scheduler() is called, it does not return immediately to caller ✓
- ☒ f. sched() and scheduler() are co-routines ✓
- ☒ g. After call to swtch() in scheduler(), the control moves to code in sched() ✓
- ☒ h. After call to swtch() in sched(), the control moves to code in scheduler() ✓

Your answer is correct.

The correct answers are: sched() and scheduler() are co-routines, When either sched() or scheduler() is called, it does not return immediately to caller, When either sched() or scheduler() is called, it results in a context switch, sched() switches to the scheduler's context, scheduler() switches to the selected process's context, After call to swtch() in scheduler(), the control moves to code in sched(), After call to swtch() in sched(), the control moves to code in scheduler(), Each call to sched() or scheduler() involves change of one stack inside swtch()

## Question 18

Partially correct

Mark 0.38 out of 0.50

The kernel ELF file contains these headers

Program Header:

```
LOAD off 0x00001000 vaddr 0x80100000 paddr 0x00100000 align 2**12
      filesz 0x00007aab memsz 0x00007aab flags r-x
LOAD off 0x00009000 vaddr 0x80108000 paddr 0x00108000 align 2**12
      filesz 0x00002516 memsz 0x0000d4a8 flags rw-
STACK off 0x00000000 vaddr 0x00000000 paddr 0x00000000 align 2**4
      filesz 0x00000000 memsz 0x00000000 flags rwx
```

mark the statemetns as True/False

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	in bootmain() the third header leads to allocation of no-memory.	✓
<input checked="" type="radio"/>	<input type="radio"/>	First header is for the code/text	✓
<input type="radio"/>	<input checked="" type="radio"/>	Third header is for stack	✗
<input checked="" type="radio"/>	<input type="radio"/>	Second header is for Data/Globals	✓

in bootmain() the third header leads to allocation of no-memory.: True

First header is for the code/text: True

Third header is for stack: True

Second header is for Data/Globals: True

## Question 19

Correct

Mark 0.50 out of 0.50

Will this code work for a spinlock() operation? The intention here is to call compare-and-swap() only if the lock is not held (the if condition checks for the same).

```
void spinlock(int *lock) {  
    {  
        while (true) {  
            if (*lock == 0) {  
                /* lock appears to be available */  
                if (!compare_and_swap(lock, 0, 1))  
                    break  
            }  
        }  
    }  
}
```

- ☐ a. No, because this breaks the atomicity requirement of compare-and-test.
- ☐ b. Yes, because there is no race to update the lock variable
- ☒ c. Yes, because no matter in which order the if-check and compare-and-swap run in multiple processes, only one process will succeed in compare-and-swap() and others will keep looping in while-loop. ✓
- ☐ d. No, because in the case of both processes succeeding in the "if" condition, both may end up acquiring the lock.

Your answer is correct.

The correct answer is: Yes, because no matter in which order the if-check and compare-and-swap run in multiple processes, only one process will succeed in compare-and-swap() and others will keep looping in while-loop.



Question **20**

Partially correct

Mark 0.36 out of 0.50

Mark statements as T/F

All statements are in the context of preventing deadlocks.

True	False		
<input checked="" type="radio"/>	<input checked="" type="radio"/>	The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order	✗
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Circular wait is avoided by enforcing a lock ordering	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	A process holding one resources and waiting for just one more resource can also be involved in a deadlock.	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Hold and wait means a thread/process holding some locks and waiting for acquiring some.	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	If a resource allocation graph contains a cycle then there is a guarantee of a deadlock	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens	✗

The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order: False

Circular wait is avoided by enforcing a lock ordering: True

Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.: True

A process holding one resources and waiting for just one more resource can also be involved in a deadlock.: True

Hold and wait means a thread/process holding some locks and waiting for acquiring some.: True

If a resource allocation graph contains a cycle then there is a guarantee of a deadlock: False

Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens: True

Question **21**

Incorrect

Mark 0.00 out of 0.50

The variable 'end' used as argument to kinit1 has the value

- ☐ a. 80000000
- ☐ b. 81000000
- ☐ c. 801154a8
- ☐ d. 80102da0
- ☐ e. 80110000
- ☒ f. 8010a48c ✗

The correct answer is: 801154a8

Question **22**

Partially correct

Mark 0.38 out of 0.50

Mark the statements as True or False, w.r.t. passing of arguments to system calls in xv6 code.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The functions like <code>argint()</code> , <code>argstr()</code> make the system call arguments available in the kernel.	✓
<input checked="" type="radio"/>	<input type="radio"/>	The arguments are accessed in the kernel code using <code>esp</code> on the trapframe.	✓
<input type="radio"/>	<input checked="" type="radio"/>	String arguments are first copied to trapframe and then from trapframe to kernel's other variables.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Integer arguments are stored in <code>eax</code> , <code>ebx</code> , <code>ecx</code> , etc. registers	✓
<input checked="" type="radio"/>	<input type="radio"/>	The arguments to system call are copied to kernel stack in <code>trapasm.S</code>	✗
<input checked="" type="radio"/>	<input type="radio"/>	Integer arguments are copied from user memory to kernel memory using <code>argint()</code>	✓
<input type="radio"/>	<input checked="" type="radio"/>	String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer	✗
<input checked="" type="radio"/>	<input type="radio"/>	The arguments to system call originally reside on process stack.	✓

The functions like `argint()`, `argstr()` make the system call arguments available in the kernel.: True

The arguments are accessed in the kernel code using `esp` on the trapframe.: True

String arguments are first copied to trapframe and then from trapframe to kernel's other variables.: False

Integer arguments are stored in `eax`, `ebx`, `ecx`, etc. registers: False

The arguments to system call are copied to kernel stack in `trapasm.S`: False

Integer arguments are copied from user memory to kernel memory using `argint()`: True

String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer: True

The arguments to system call originally reside on process stack.: True

Question **23**

Correct

Mark 0.50 out of 0.50

The variable `$stack` in `entry.S` is

- ☐ a. located at less than `0x7c00`
- ☐ b. located at the value given by `%esp` as setup by `bootmain()`
- ☐ c. located at 0
- ☒ d. a memory region allocated as a part of `entry.S` ✓
- ☐ e. located at `0x7c00`

The correct answer is: a memory region allocated as a part of `entry.S`

Question **24**

Correct

Mark 0.50 out of 0.50

In the code below assume that each function can be executed concurrently by many threads/processes. Ignore syntactical issues, and focus on the semantics.

This program is an example of

```
spinlock a, b; // assume initialized
thread1() {
    spinlock(a);
    //some code;
    spinlock(b);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
thread2() {
    spinlock(a);
    //some code;
    spinlock(b);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
```

- ☐ a. Livelock
- ☐ b. Deadlock
- ☐ c. Deadlock or livelock depending on actual race
- ☐ d. Self Deadlock
- ☒ e. None of these ✓

Your answer is correct.

The correct answer is: None of these

Question **25**

Partially correct

Mark 0.25 out of 0.50

Which of the following is not a task of the code of swtch() function

- ☒ a. Jump to next context EIP ✗
- ☒ b. Change the kernel stack location ✓
- ☐ c. Load the new context
- ☐ d. Save the old context
- ☐ e. Switch stacks
- ☐ f. Save the return value of the old context code

The correct answers are: Save the return value of the old context code, Change the kernel stack location

Question **26**

Correct

Mark 0.50 out of 0.50

Why V2P\_WO is used in entry.S and not V2P ?

- ☐ a. The two macros are different. They lead to different calculations.
- ☐ b. It's a mistake. They could have used the same macro in both places.
- ☐ c. The typecasting has the effect of creating virtual address, while without typecast we get physical address.
- ☒ d. Because entry.S is an assembly code file and assemblers do not know about data types and type casting. ✓
- ☐ e. Because the processor can not do a type casting at run time

Your answer is correct.

The correct answer is: Because entry.S is an assembly code file and assemblers do not know about data types and type casting.

Question **27**

Partially correct

Mark 0.08 out of 0.50

Which of the following is done by mappages()?

- ☐ a. allocate page directory if required
- ☒ b. create page table mappings to the range given by "pa" and "pa + size" ✓
- ☐ c. allocate page table if required
- ☒ d. allocate page frame if required ✗
- ☒ e. create page table mappings for the range given by "va" and "va + size" ✓

The correct answers are: create page table mappings for the range given by "va" and "va + size", allocate page table if required, create page table mappings to the range given by "pa" and "pa + size"

Question **28**

Correct

Mark 0.50 out of 0.50

In an ext2 file system, if the block size is 4KB and partition size is 64 GB, then the number of block groups will be:

Answer:  ✓

size \* 1024 \* 1024 / 4 --> no of blocks

each group = 8 \* 4 \* 1024 blocks = 32768 blocks

so size \* 1024 \* 1024 / (4 \* 32768) number of groups

The correct answer is: 512.00

## Question 29

Correct

Mark 0.50 out of 0.50

Mark statements as True/False w.r.t. the creation of free page list in xv6.

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	if(kmem.use_lock) acquire(&kmem.lock); this "if" condition is true, when kinit2() runs because multi-processor support has been enabled by now.	✓ No. kinit2() calls kfree() and then initializes use_lock.
<input checked="" type="radio"/>	<input type="radio"/>	The pointers that link the pages together are in the first 4 bytes of the pages themselves	✓
<input checked="" type="radio"/>	<input type="radio"/>	the kmem.lock is used by kfree() and kalloc() only.	✓
<input checked="" type="radio"/>	<input type="radio"/>	if(kmem.use_lock) acquire(&kmem.lock); is not done when called from kinit1() because there is no need to take the lock when kinit1() is running because interrupts are disabled and only one processor is running	✓
<input checked="" type="radio"/>	<input type="radio"/>	kmem.use_lock is set to 1 after free page list is created, so that kmem.lock is taken before accessing kmem.freelist.	✓
<input type="radio"/>	<input checked="" type="radio"/>	free page list is a singly circular linked list.	✓ it's singly linked NULL terminated list.

```
if(kmem.use_lock)
```

```
    acquire(&kmem.lock);
```

this "if" condition is true, when kinit2() runs because multi-processor support has been enabled by now.: False

The pointers that link the pages together are in the first 4 bytes of the pages themselves: True

the kmem.lock is used by kfree() and kalloc() only.: True

```
if(kmem.use_lock)
```

```
    acquire(&kmem.lock);
```

is not done when called from kinit1() because there is no need to take the lock when kinit1() is running because interrupts are disabled and only one processor is running: True

kmem.use\_lock is set to 1 after free page list is created, so that kmem.lock is taken before accessing kmem.freelist.: True

free page list is a singly circular linked list.: False

## Question 30

Correct

Mark 0.50 out of 0.50

Doing a lookup on the pathname /a/b/b/c/d for opening the file "d" requires reading  ✓ no. of inodes. Assume that there are no hard/soft links on the path.

Write the answer as a number.

The correct answer is: 6

◀ Quiz-1 (15 Marks)

Jump to...

ESE(60 Marks) ▶